



Future Network Architectures **Recursive Internet Architecture (RINA)**


Dimitri Staessens – Ghent Uni. iMinds (BE)

Francesco Salvestrini – Nextworks s.r.l. (IT)


Miquel Tarzan – Fundació i2CAT (ES)

Challenges faced by (Inter)network engineers

- explosion in the complexity of the overall system (hundreds of protocols and thousands of standards documents)
- weak security
- scalability issues with the routing system
 - (IPv6/BGP multihoming)
 - Mobile end-users
- no QoS support



Prefixes announced on the Internet



The Internet is a live environment

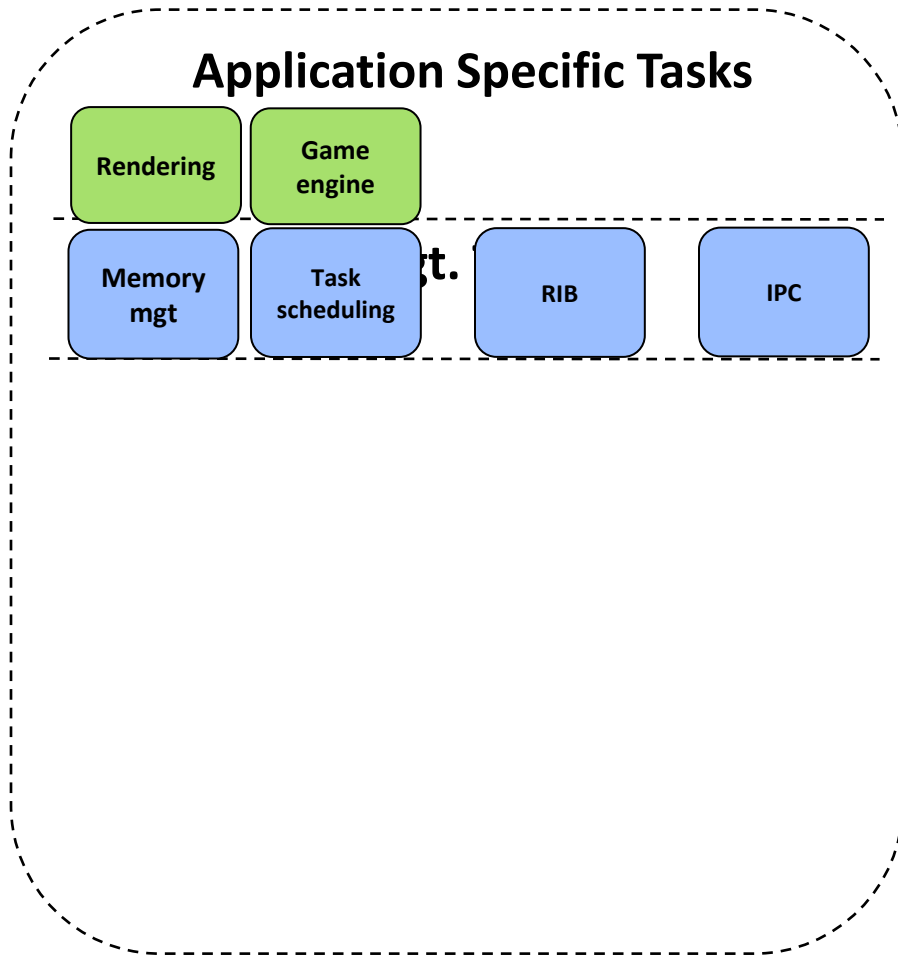
- ever growing customer base
- ever growing number of devices
- new and more demanding services
- RAD of services
- fast deployment
- “whac-a-mole” approach to problemsolving



A brief introduction to the Recursive Internet Architecture

RINA

Application Process



Components

- Application specific tasks
- Management tasks


Mechanism

- Static, invariant parts

Policy

- Dynamic, variant parts
- occurs in pairs
 - Sender
 - Receiver

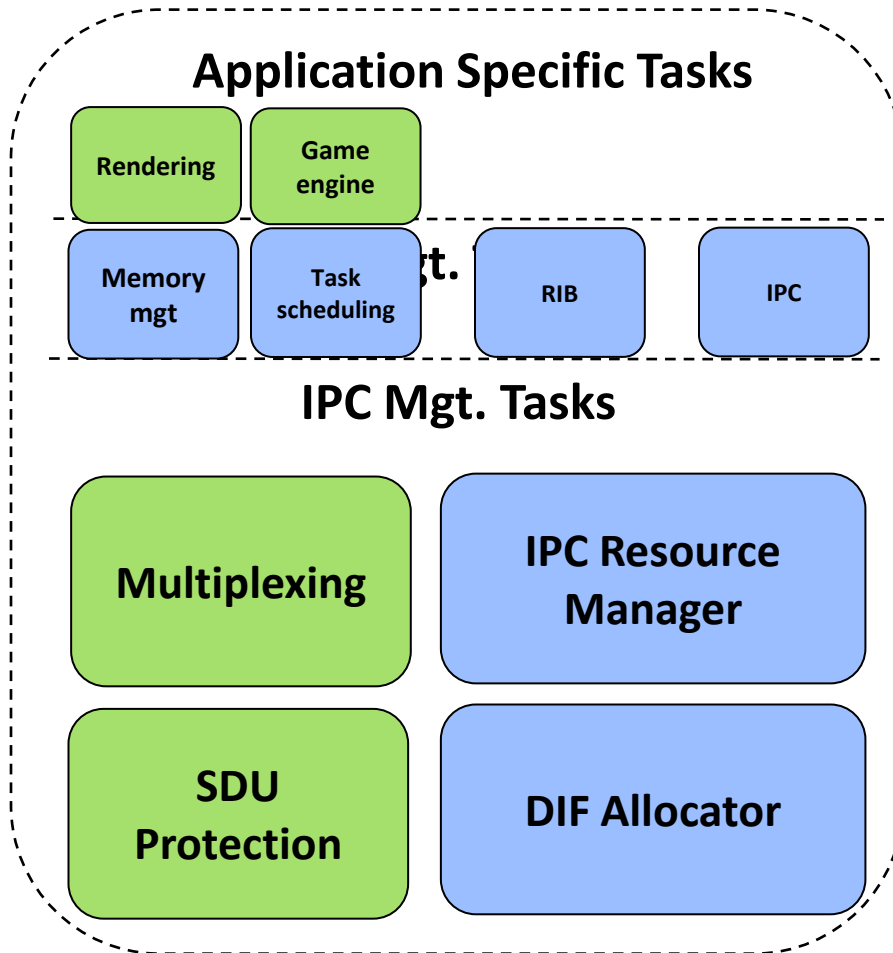
Distributed Application



Processing system: hardware and software capable of executing programs as Application Processes that can coordinate via shared memory (“test and set”)

Computing system: a collection of processing systems under the same management domain with no restrictions on connectivity

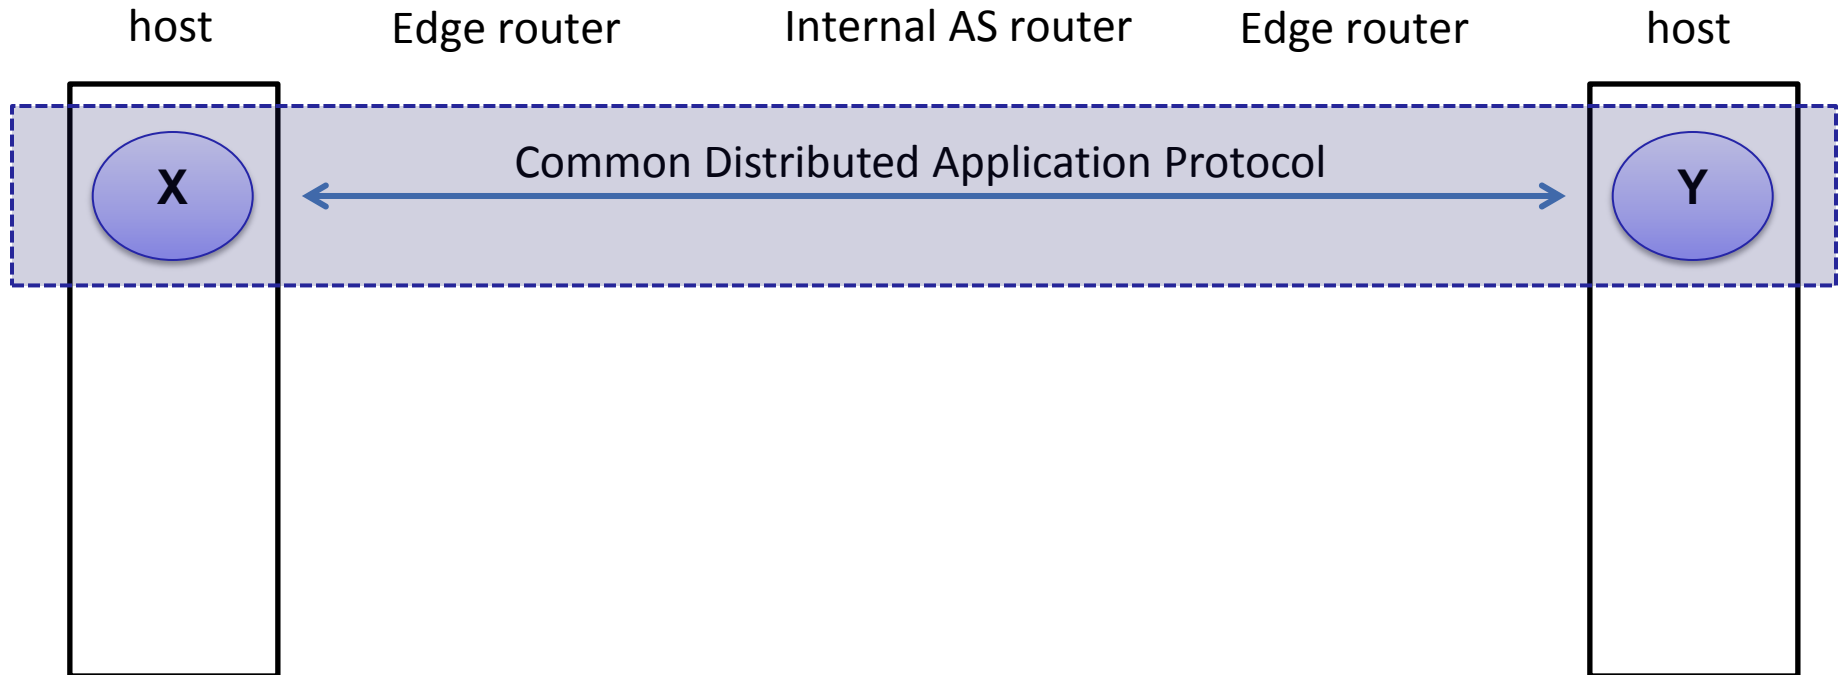
Application Process



Components

- Application specific tasks
- Management tasks
- IPC Management tasks
 - DIF Allocator: Finds remote application processes
 - IRM: manages DA requests
 - Multiplexing: SDU's from different tasks
 - SDU protection: Integrity and security

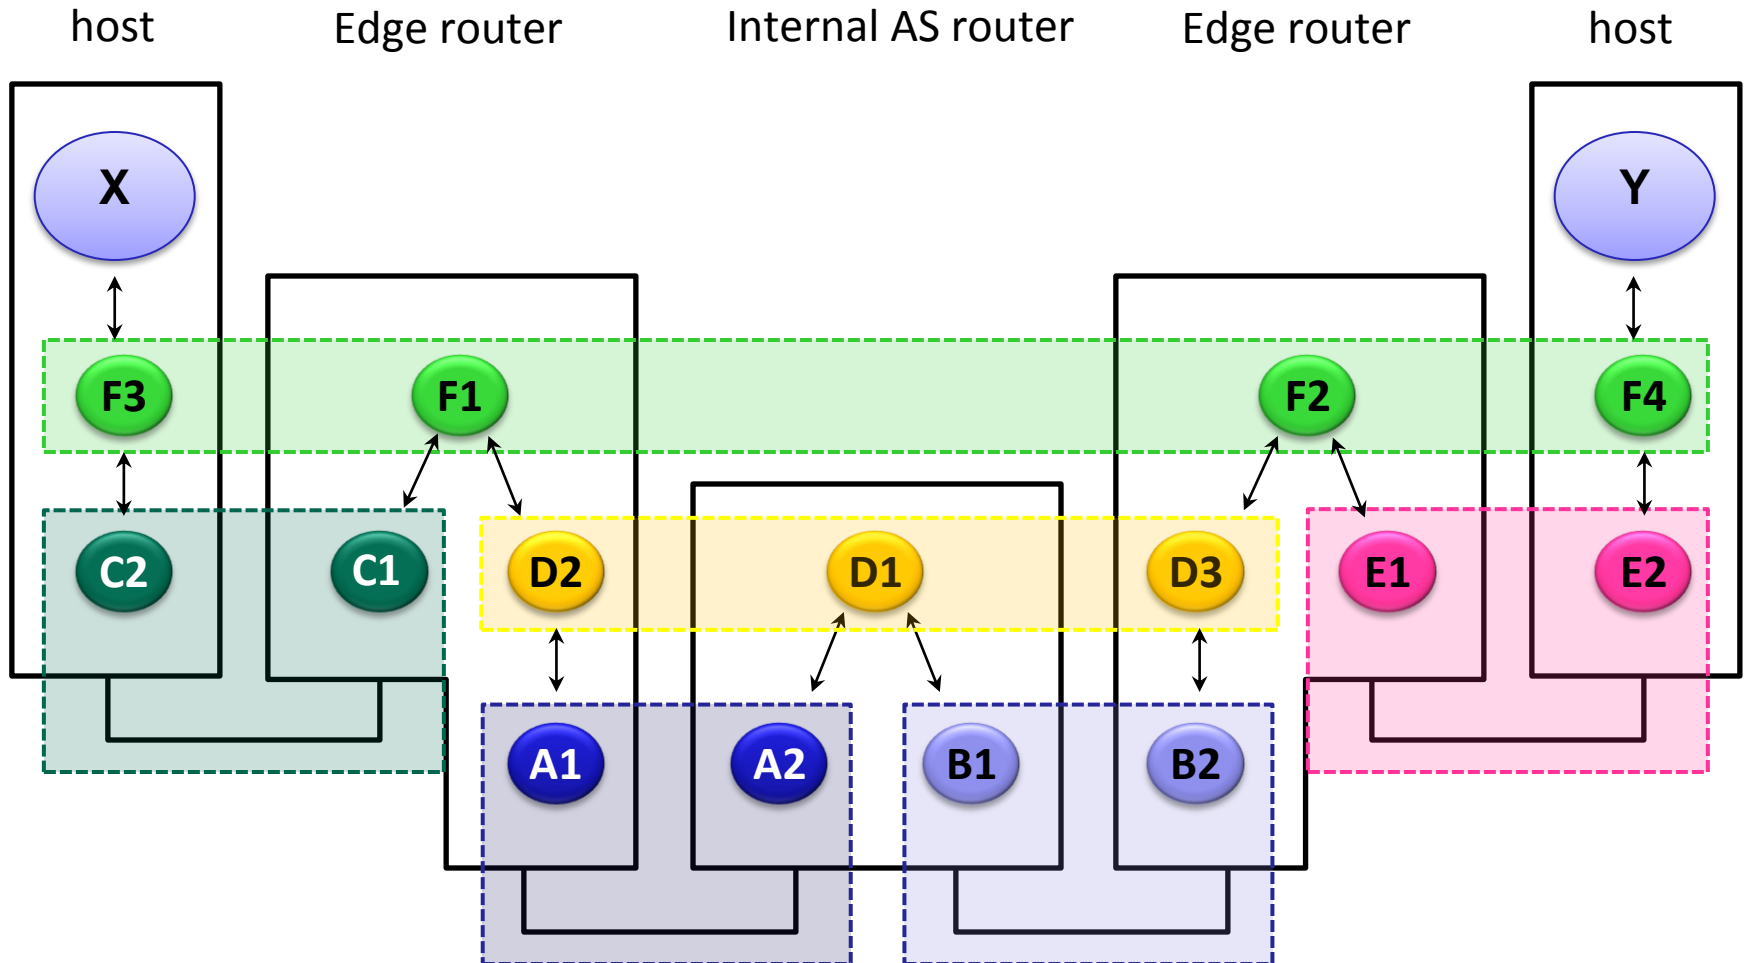
Distributed Applications Provide IPC services






Common Distributed Application Protocol (CDAP)

- Perform operations on RIB objects
 - **Create/Delete**
 - **Read/Write**
 - **Start/Stop**
- But what about different applications?
 - The objects they manipulate
 - Control and sequencing of operations
 - ...

Distributed Applications Provide IPC services

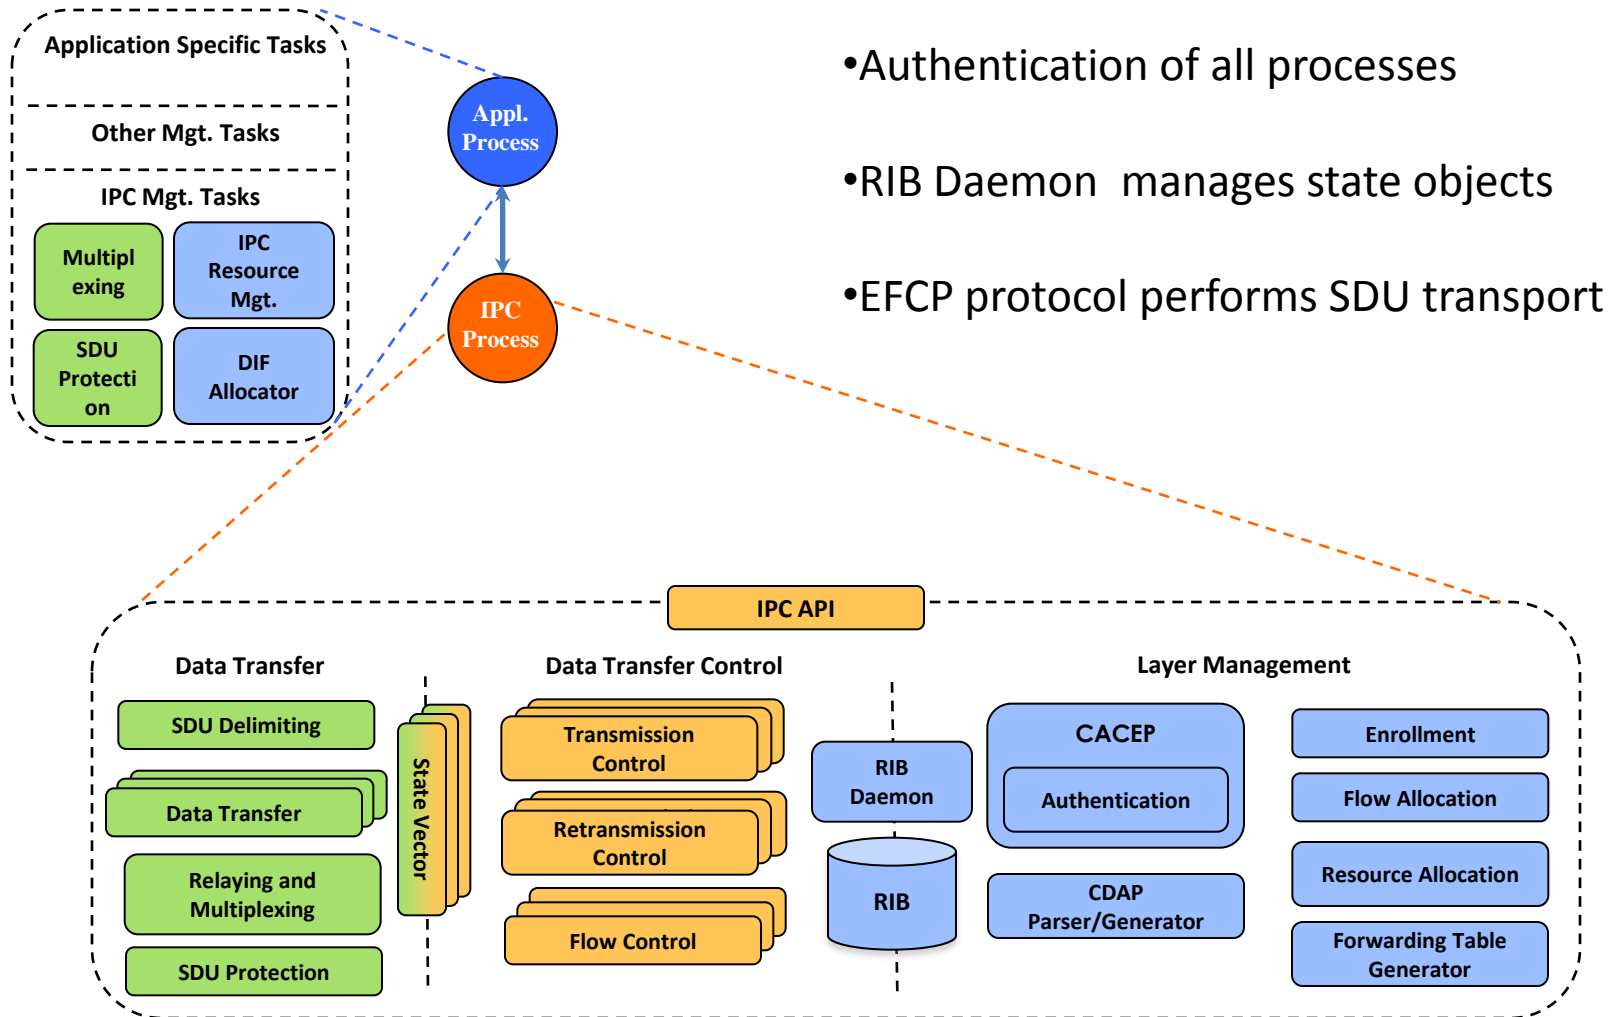


Effectively extending the IPC model



- APs communicate using a port, identified by a portId
- 6 operations:
 - `int _registerApp(appName, List<difName>)`
 - `portId _allocateFlow(destAppName, List<QoSParams>)`
 - `int _write(portId, sdu)`
 - `sdu _read(portId)`
 - `int _deallocate(portId)`
 - `int _unregisterApp(appName, List<difName>)`
- QoSParams are defined in a technology-agnostic way
 - Bandwidth-related, delay, jitter, in-order-delivery, loss rates, ...

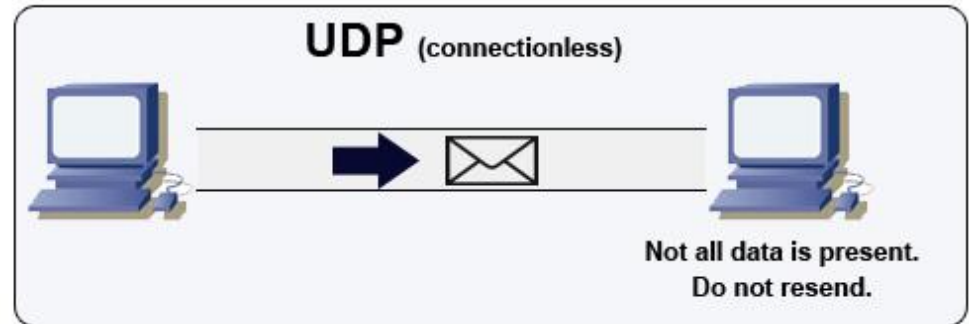
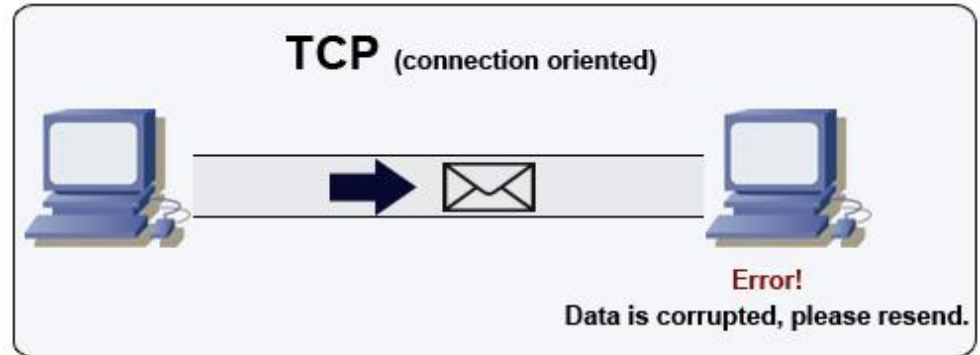
The IPC process



- Authentication of all processes
- RIB Daemon manages state objects
- EFCP protocol performs SDU transport

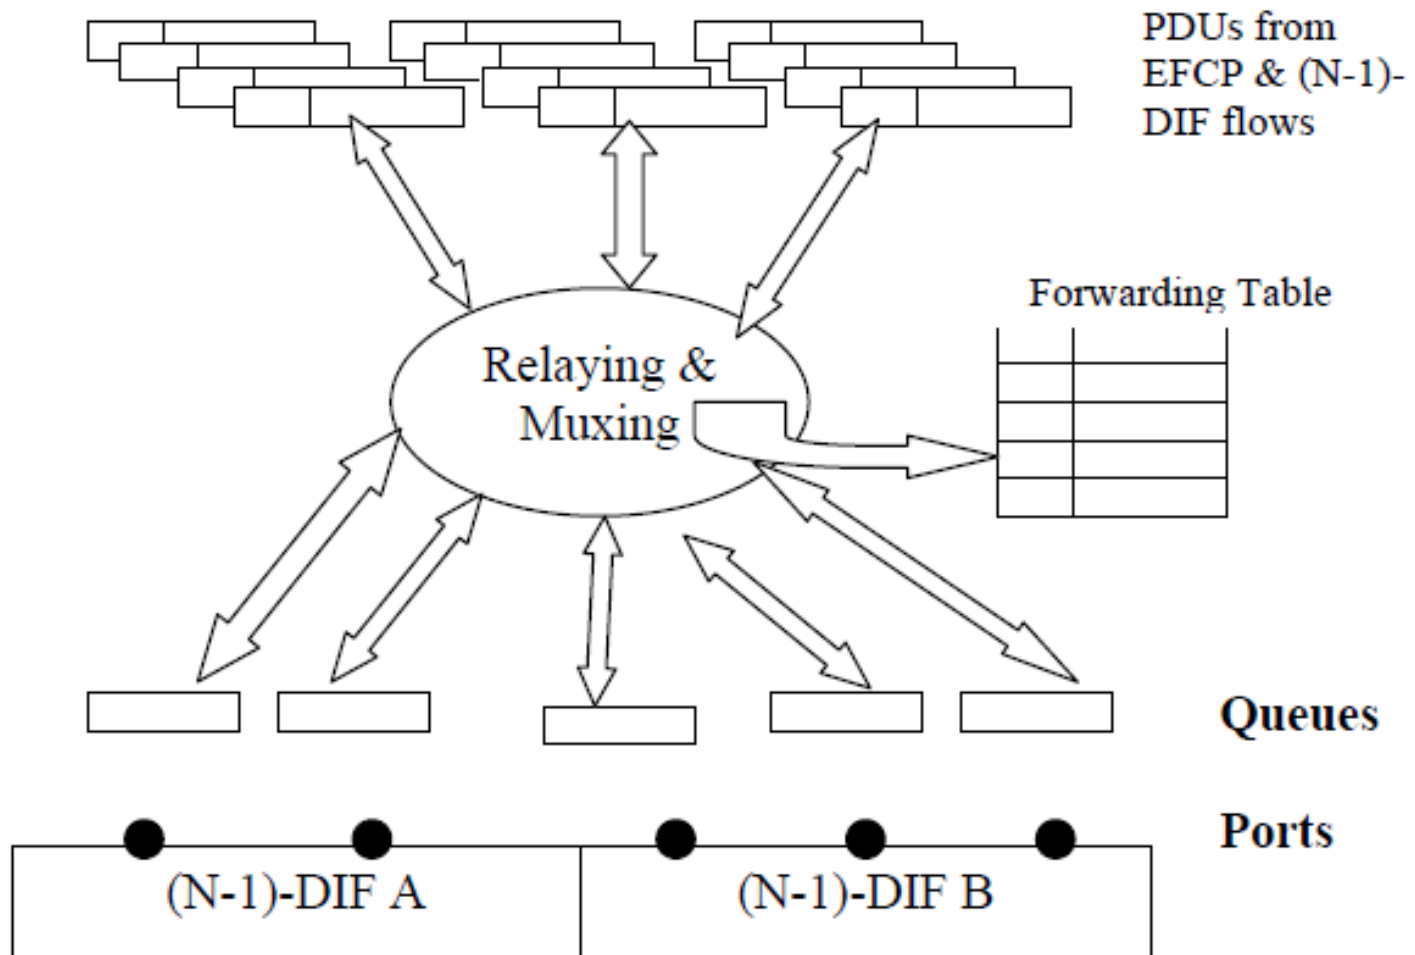
Error and Flow Control Protocol

- DTP
 - Fragmentation
 - Reassembly
 - Sequencing
 - Concatenation
 - Separation
- DTCP
 - Transmission control
 - Retransmission control
 - Flow control
- Loosely coupled by a state vector
- Based on Delta-t




Delta-t (Watson, 1981)

- Developed at L.Livermore labs, unique approach.
 - Assumes all connections exist all the time.
 - keep caches of state on ones with recent activity
- Watson proves that the conditions for distributed synchronization are met if and only if 3 timers are bounded:
 - Maximum Packet Lifetime: **MPL**
 - Maximum number of Retries: **R**
 - Maximum time before Ack: **A**
- That no explicit state synchronization, i.e. hard state, is necessary.
 - SYNs, FINs are unnecessary
- 1981:Watson shows that TCP has all three timers and more.

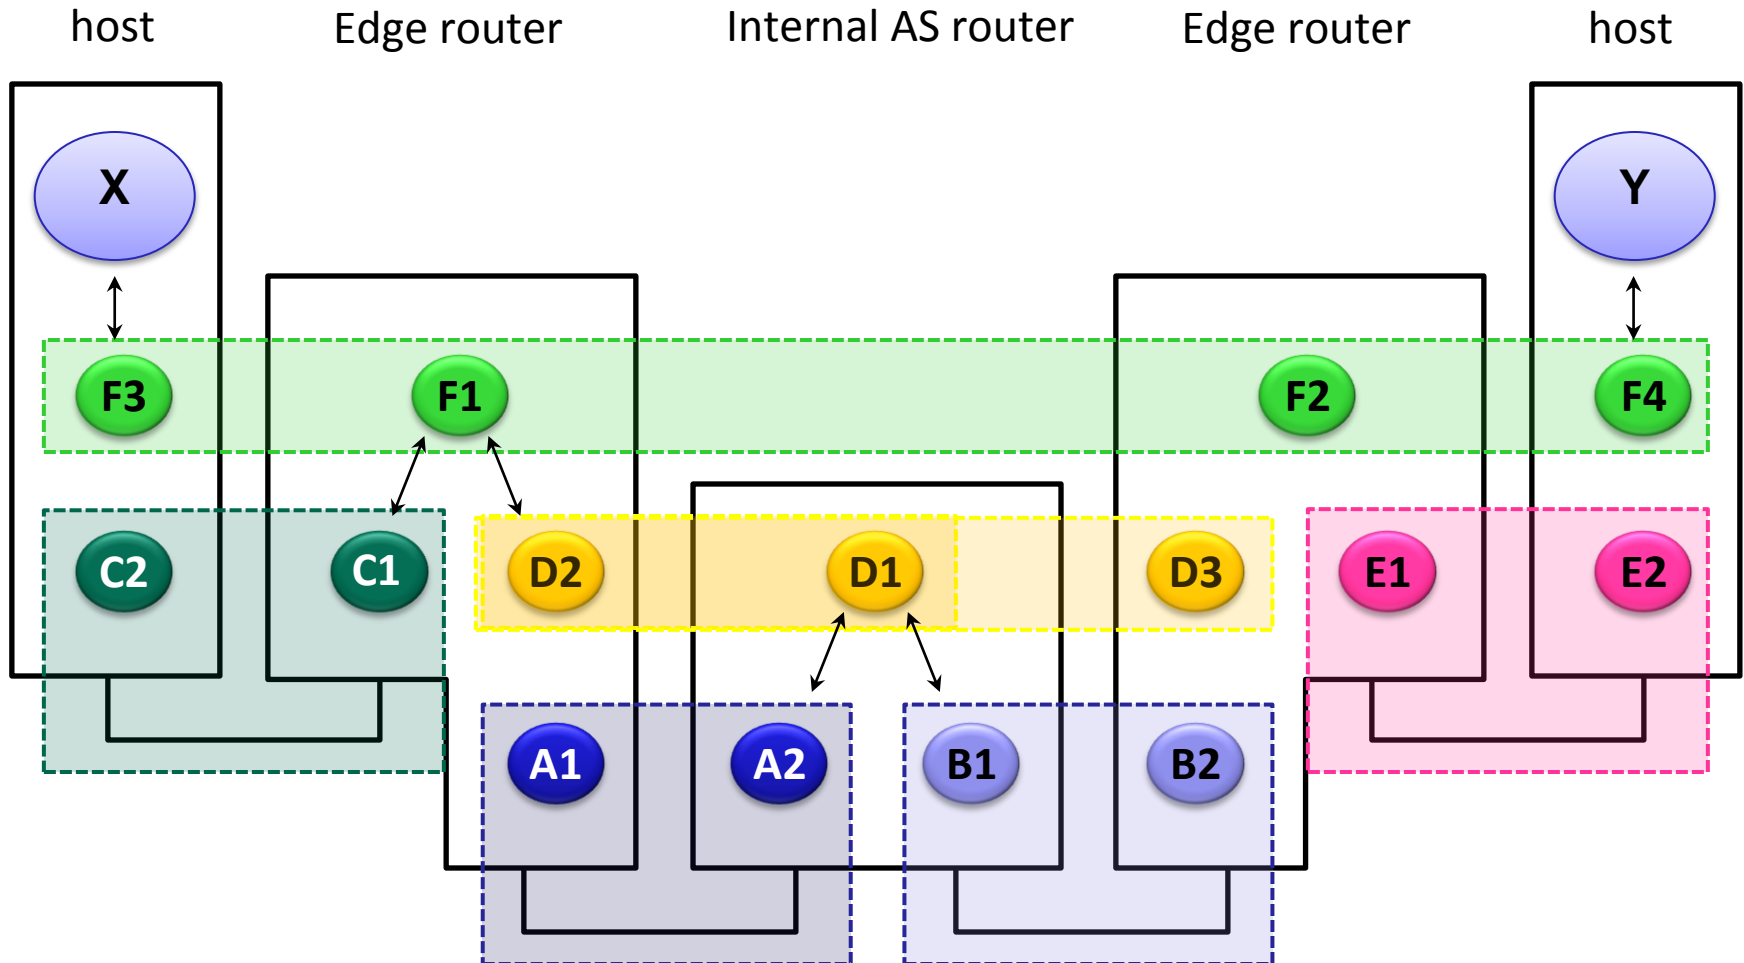


- Wrap a technology with the IPC API
 - Physical medium
 - Legacy technology
 - Ethernet
 - IP
 - Hypervisors
- Not required to add functionality
- So it's an “incomplete” DIF

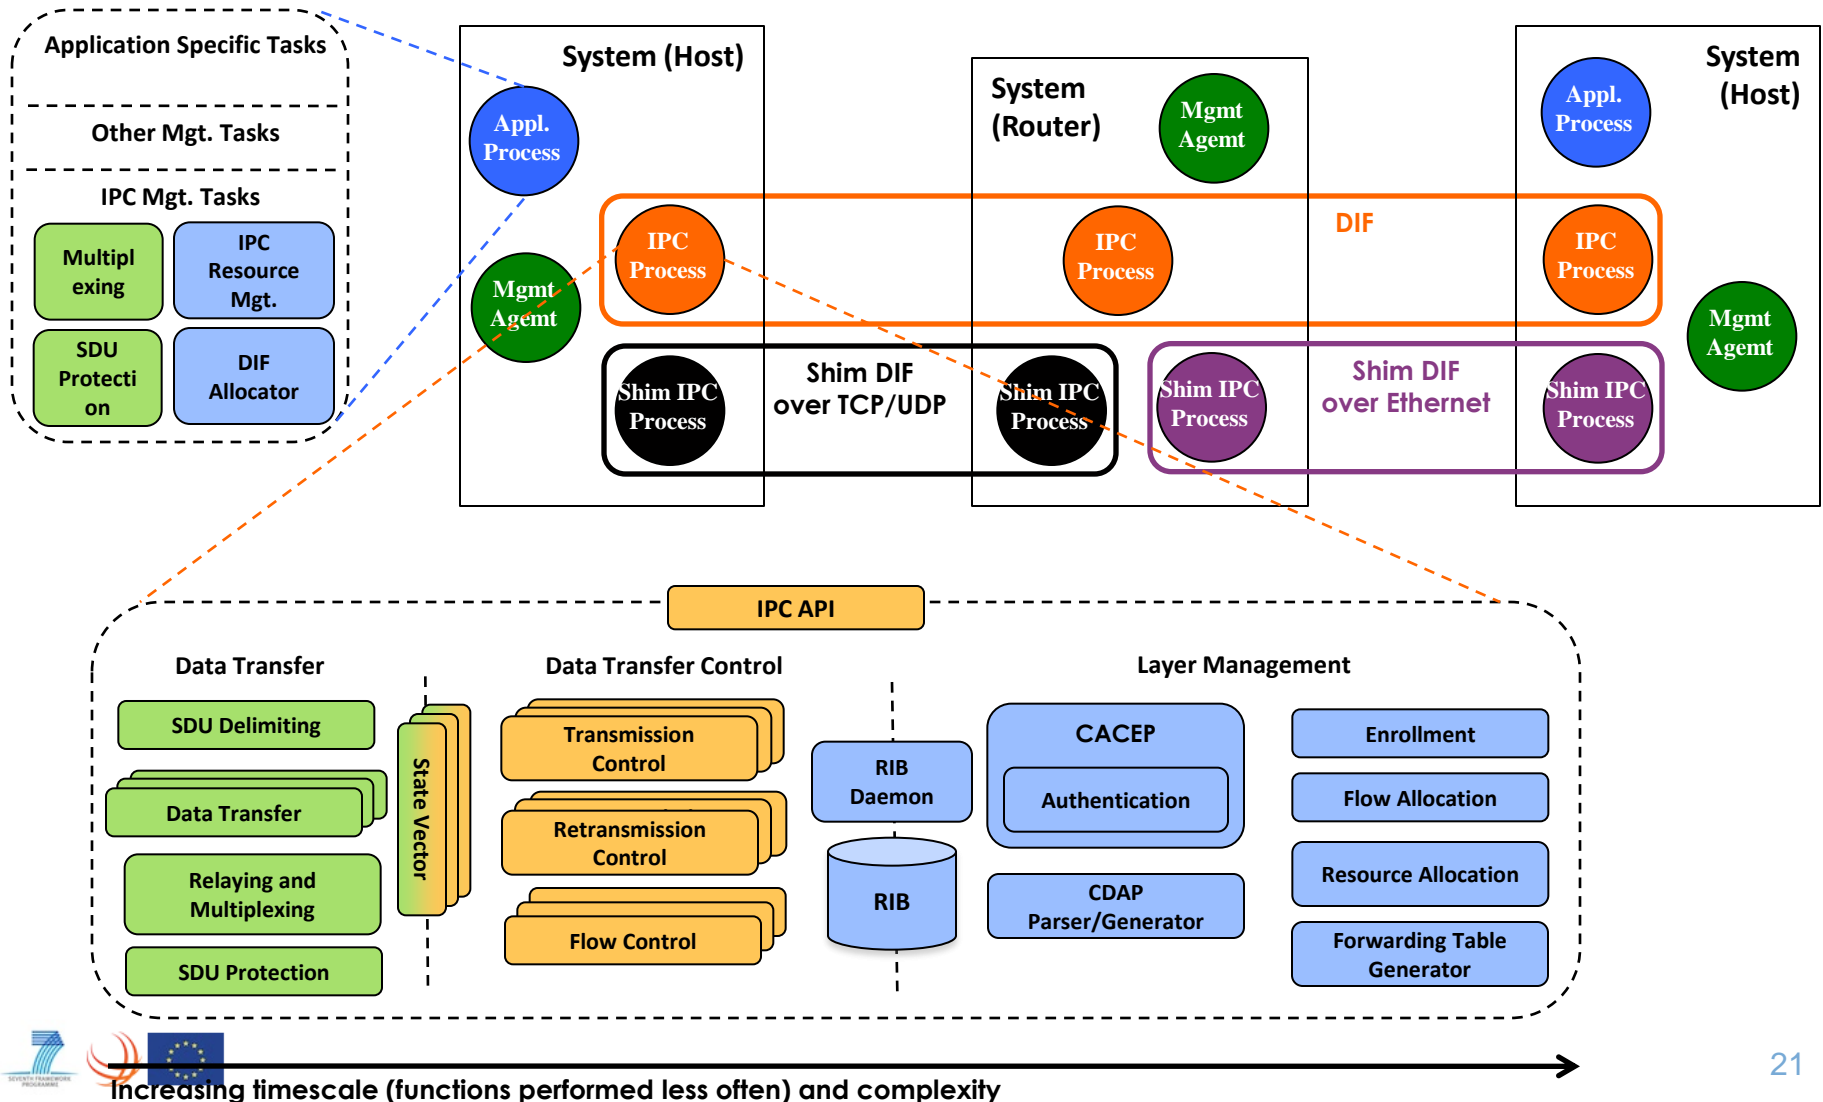
Basic concept of RINA



Bootstrapping a RINA network



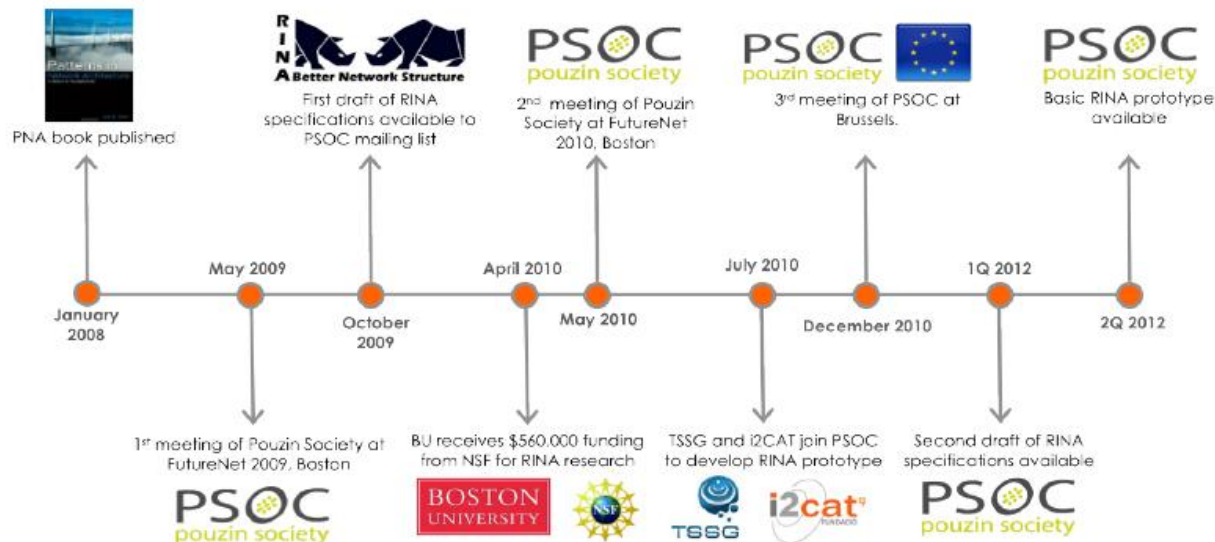
Architectural Model



PROTOTYPES

Pre-existing prototypes

- The first RINA prototype dates back in 2012
- This first implementation was a joint development of BU, i2CAT and WIT-TSSG
 - Targeting the validation of the theory and specs
 - Java based, user-space



- TRIA LLC (US) built another (closed-source) prototype (C based, user-space)
- **Later:**
 - EC funded FP7 IRATI and PRISTINE projects to advance the research on RINA
 - GEANT showed interest in RINA as well, funding the IRINA project

FP7 IRATI – OVERVIEW

IRATI - Introduction


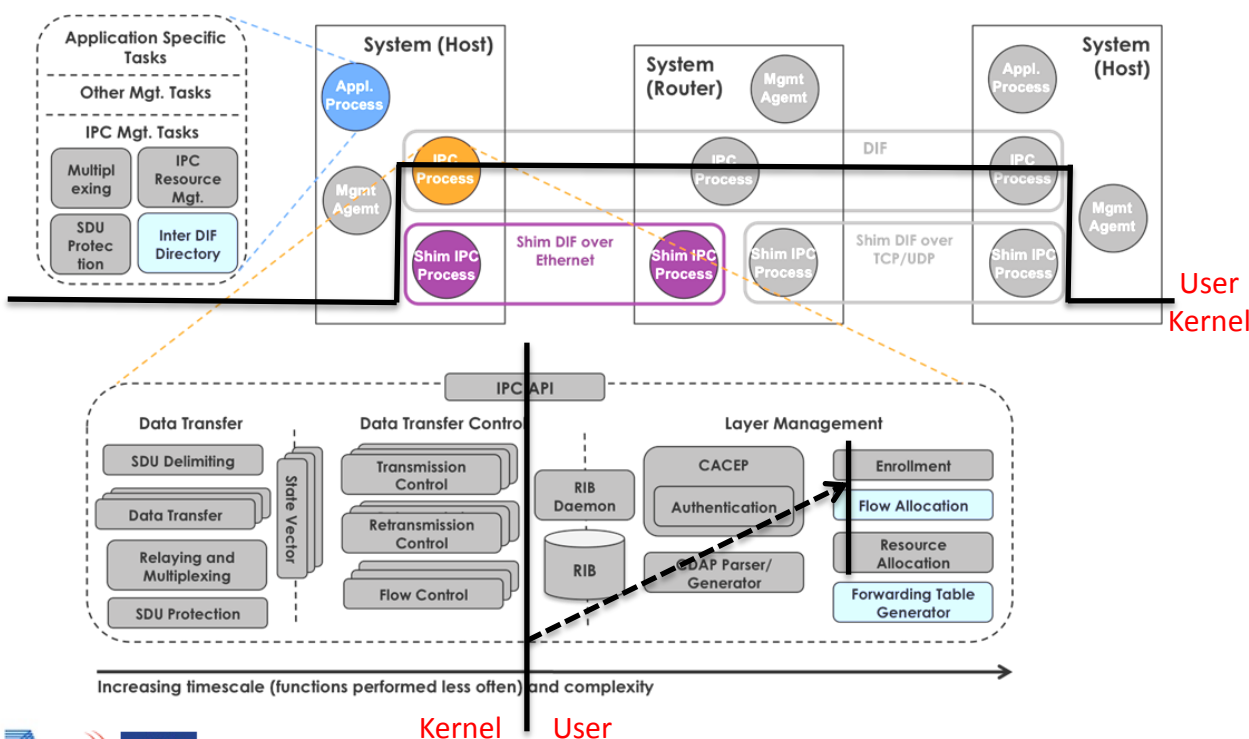
- FP7 Project – **Jan 2013** to **Dec 2014** (2 years)
- 4 partners
 - [Research] Fundació Privada i2CAT (Spain)
 - [Research] iMinds VZW (Belgium)
 - [SME] Nextworks s.r.l. (Italy)
 - [Industry] Interoute (UK/Italy)
 - [Academia] Boston University (US)



- IRATI' objectives:
 1. Enhancement of the RINA architecture reference model and specifications, focusing on DIFs over Ethernet
 2. **RINA open source prototype over Ethernet for a UNIX-like OS**
 3. Experimental validation of RINA and comparison against TCP/IP
 4. **RINA prototype for Hypervisors**
 5. **Interoperability with the PSOC RINA prototype over UDP/IP**
 6. Provide feedback to OFELIA in regards to the prototyping of a clean slate architecture
- The project targets the design and implementation of core functionalities **at processing system level**:
 - IPC Process / IPC Manager daemons
 - Transport and management layer
- Software-wise, the project:
 - Built a RINA SW prototype from scratch [**ready**]
 - Release it as FOSS [**end of October**]
 - tries to build up an Open Source community around the prototype

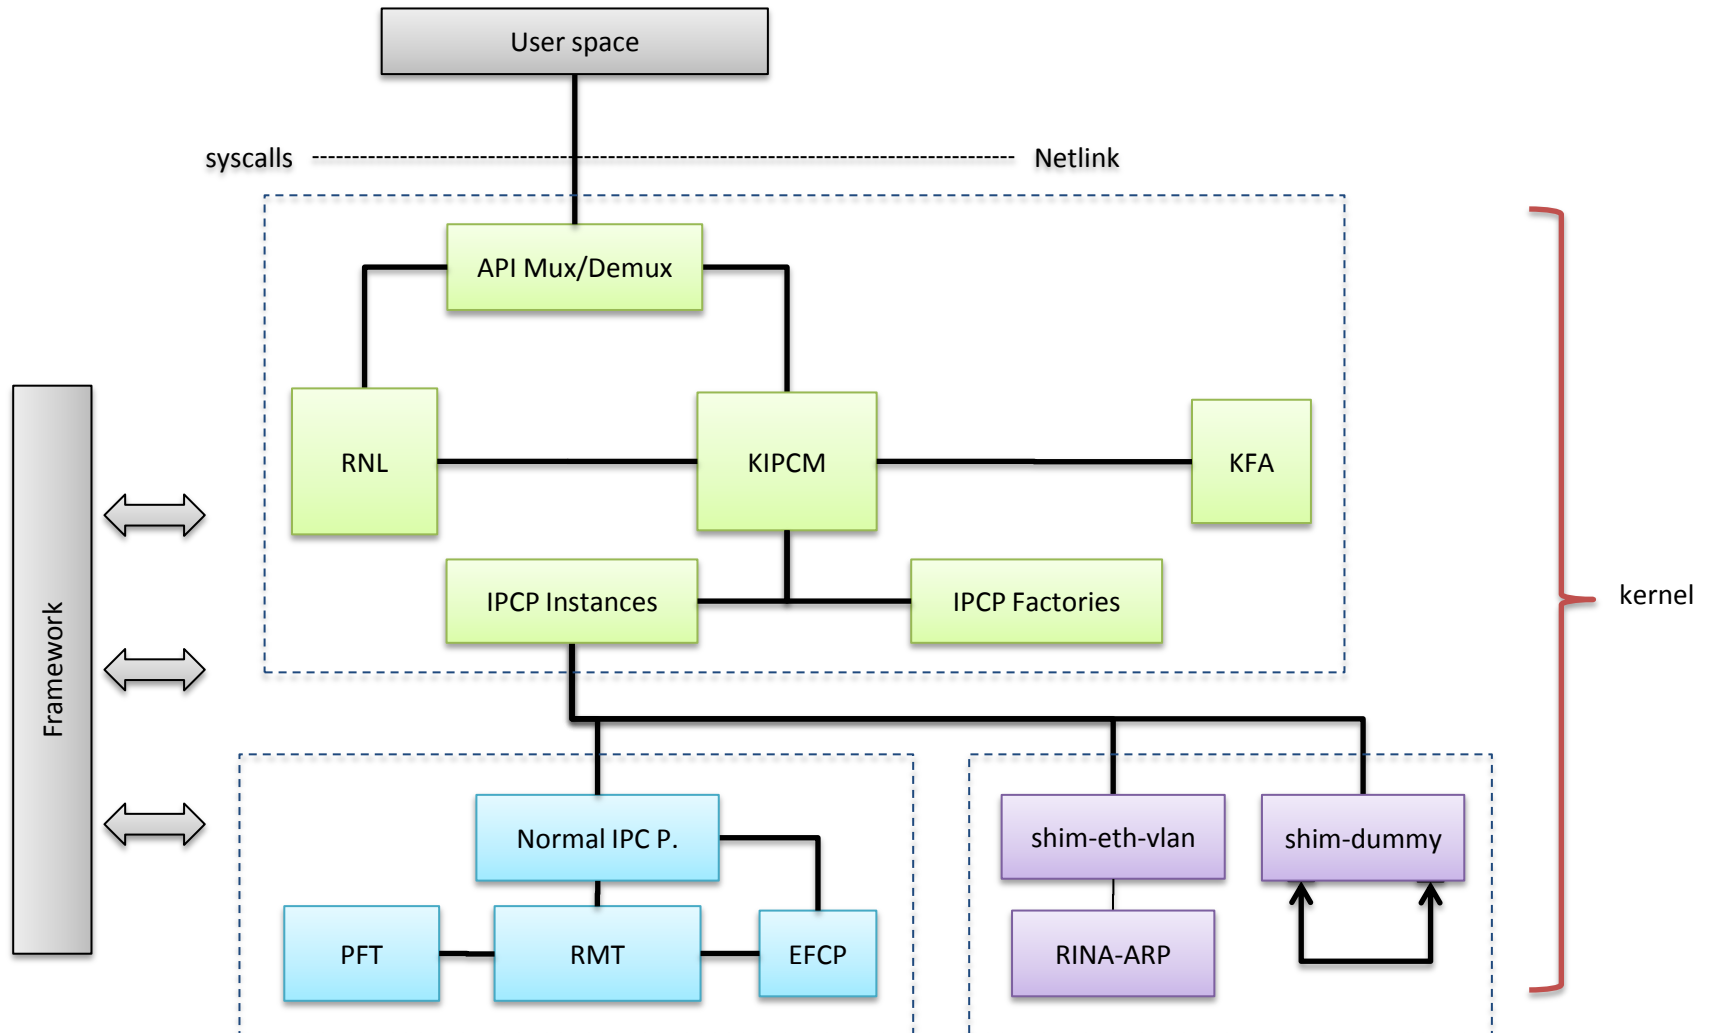
IRATI - Design decisions & fast/slow paths

- We split the RINA architecture in different “lanes”
 - Stringent timings** → **Fast-path** → **kernel-space** [TRANSPORT LAYERS]
 - loose timings** → **Slow-path** → **user-space** [MANAGEMENT LAYERS]
- Placing SW components on different lanes, **depending on their timing requirements**



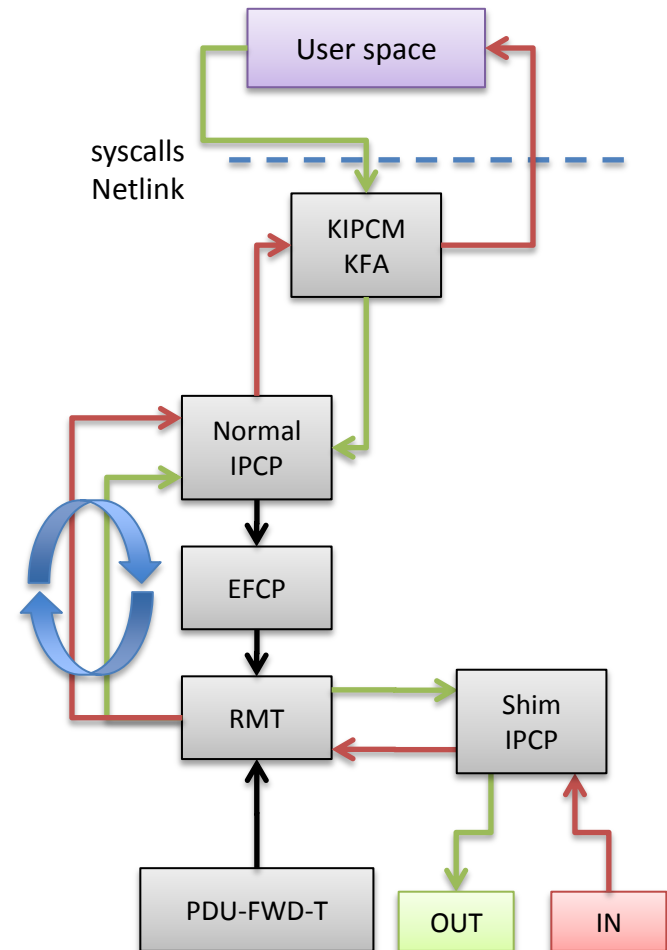
FP7 IRATI – KERNEL SPACE (THE TRANSPORT LAYERS)

The kernel-space HL SW arch



KIPCM & KFA

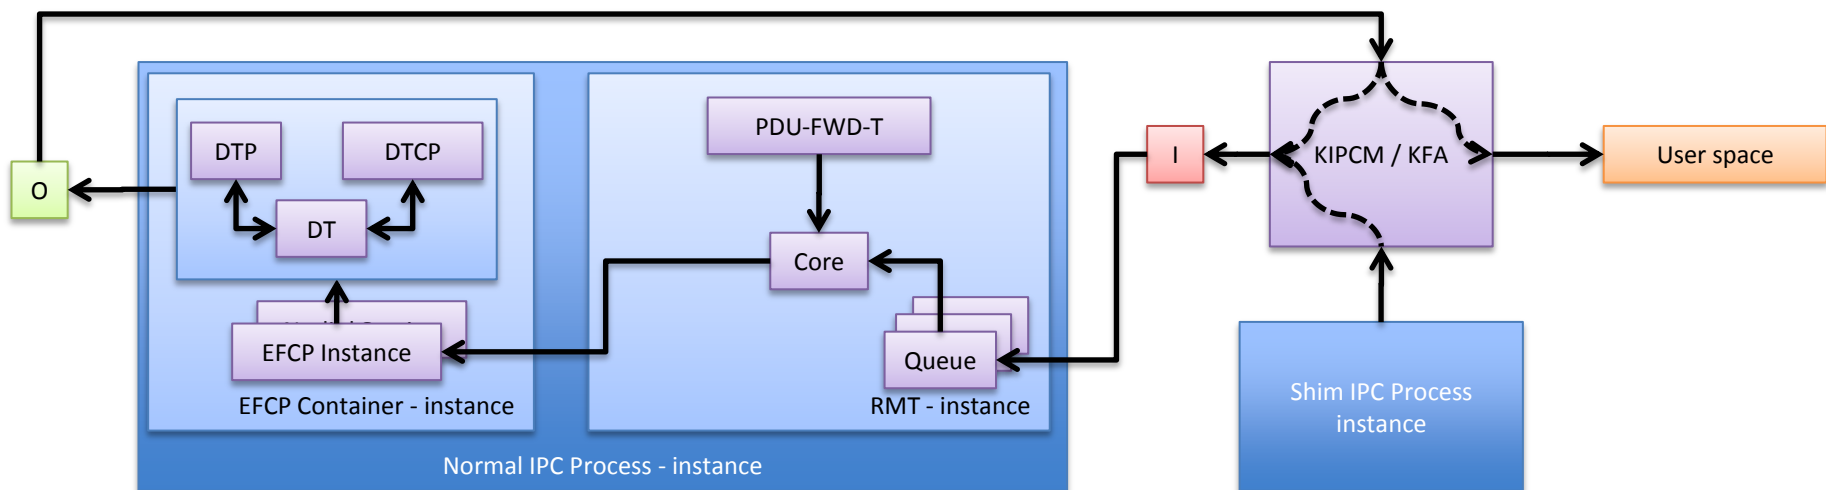
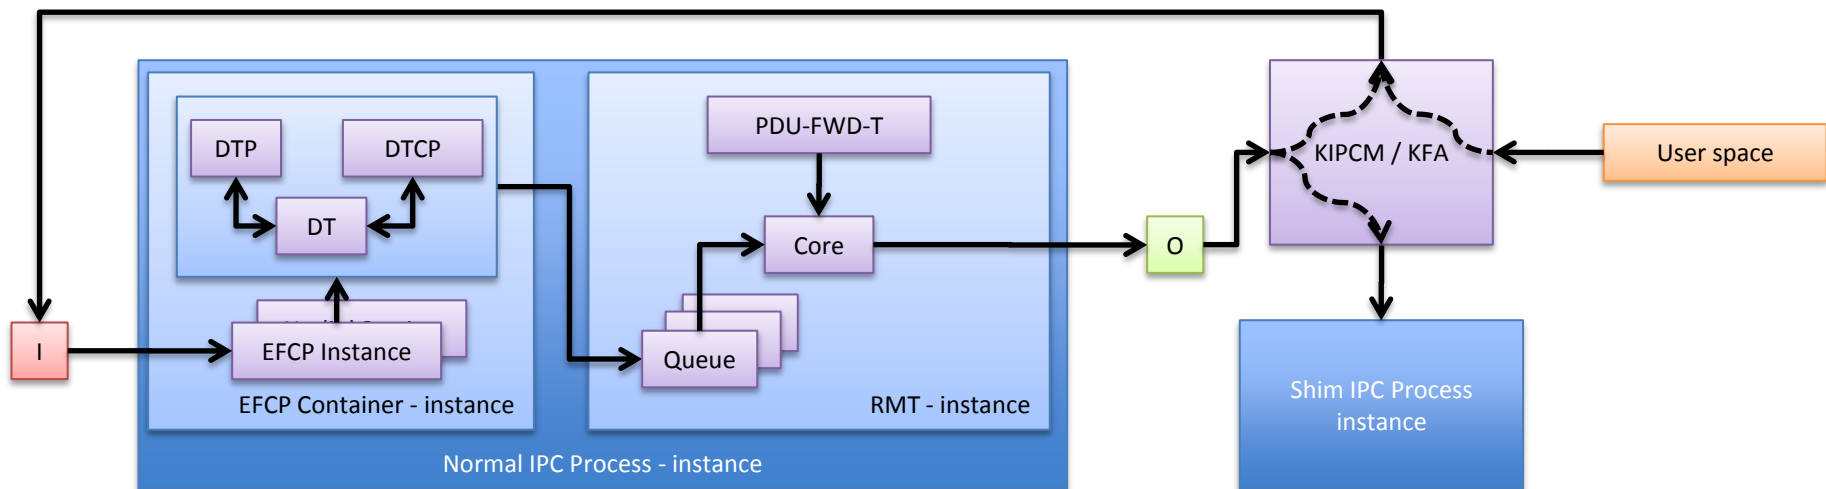
- **The KIPCM**
 - **Manages the lifecycle the IPC Processes and KFA**
 - **Abstract IPC Process instances**
 - Same API for all the IPC Processes
 - Regardless the type
 - maps: ipc-process-id → ipc-process-instance
- **The KFA**
 - **Manages ports and flows**
 - **Ports**
 - Flow handler
 - Port ID Manager
 - **Flows**
 - maps: port-id → ipc-process-instance
- **KIPCM + KFA**
 - Decouple user-interface from IPC Processes
 - **Are the Initial point where “recursion” is transformed into “iteration”**



- **There are two “major” types of IPC Processes:**
 - Normal (EFCP + RMT ...)
 - Shims
- **The interface is the same** regardless of their type
- **Each IPC Process implements its “core” details:**
 - Normal IPC Processes:
 - **The stack provides the implementation for all of them**
 - Shim IPC Processes:
 - **Each Shim IPC Processes provide its implementation**
- IPC Process instances are created via “templates”, instantiated by IPC Process Factories (OOD/OOP)

- The shims are the “lowest” components in the stack
- They have one interface:
 - **North-Bound:** The IPC API (as all the other IPC Processes)
- They wrap the underlying technology (e.g. 802.1Q)
- There are currently 4 shims available:
 - **shim-eth-vlan:**
 - As defined in the spec, runs over 802.1Q
 - **shim-hv:**
 - Targets hypervisor-based environments (KVM/Qemu and Xen)
 - Allows removing unnecessary layering commonly used in traditional VM/HV environments (e.g. bridges, virtual-NICs), such layers ease the adoption BUT:
 - Reduce performances
 - Increase maintenance costs
 - **shim-tcp-udp**
 - Targets RINA over TCP/UDP
 - **shim-dummy:**
 - Not a “real” shim, it's for debugging/testing purposes
 - It's a sort of “loopback” shim (i.e. confined into a single host)

(S | P)DUs workflows

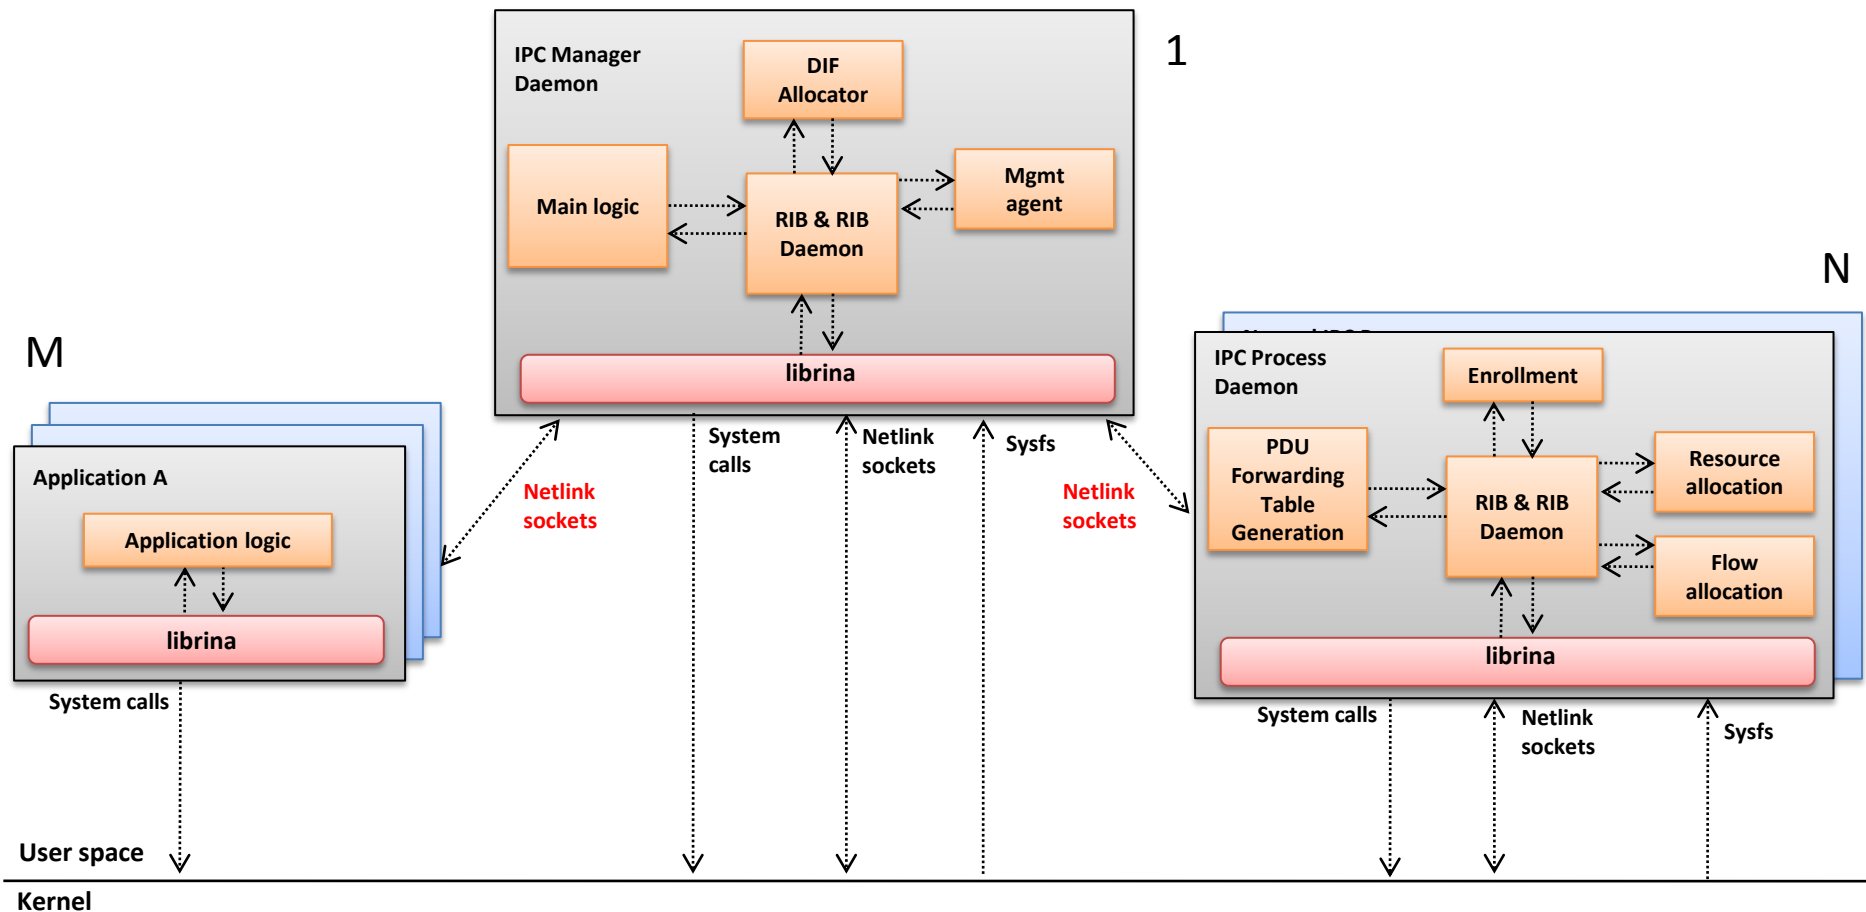


FP7 IRATI – USER SPACE (LAYER MANAGEMENT)


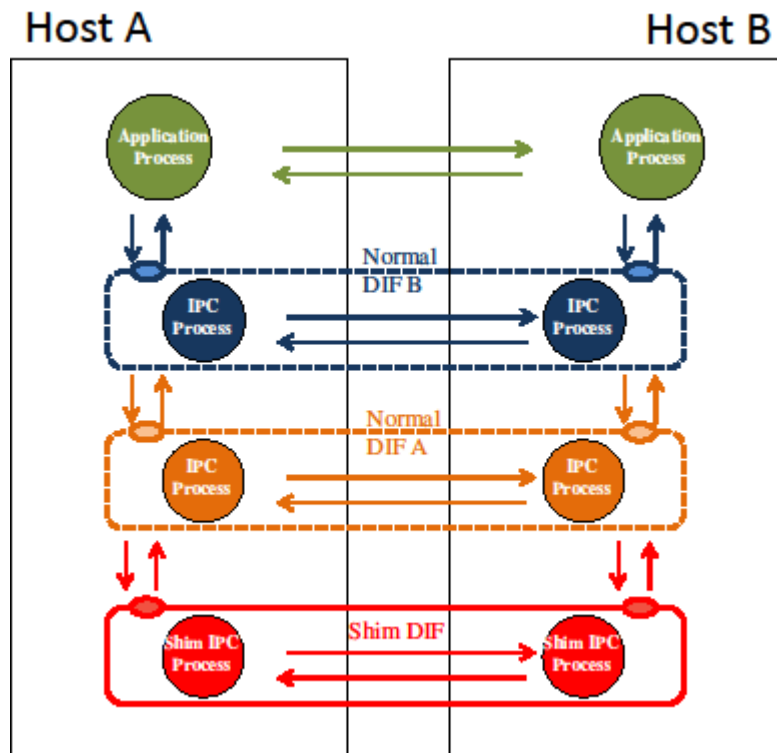
IPC Process & IPC Manager daemons

- **IPC Manager Daemon (as an OS process)**
 - **Manages the IPC Processes lifecycle**
 - **Broker between applications and IPC Processes**
 - **Local management agent**
 - **DIF Allocator client** (to search for applications not available through local DIFs)
- **IPC Process Daemon (as an OS process)**
 - **Layer Management components of the IPC Process**
 - RIB Daemon, RIB
 - CDAP parsers/generators
 - CACEP
 - Enrollment
 - Flow Allocation
 - Resource Allocation
 - PDU Forwarding Table Generation
 - Security Management
- **These daemons:**
 - **Run as separate OS processes**
 - **Rely on a common framework (librina)**

User space HL SW arch



Prototype' performances





- Theoretical maximum
- Shim IPC process for 802.1Q
- Normal IPC process over the shim IPC process for 802.1Q
- .-.-.- Normal IPC process over a normal IPC process over the shim IPC process for 802.1Q

FP7 PRISTINE – OVERVIEW

PRISTINE - Intro



- FP7 Project
 - Starts **Jan 2014**, ends **Dec 2016** (3 years)
 - 15 Partners (**Research**, **SMEs** and **Industry**)



PRISTINE - Objectives

- Design and implement the innovative internals of the RINA architecture that include the **programmable functions** for:
 - **security** of content and application processes,
 - supporting **QoS and congestion control** in aggregated levels, providing **protection and resilience**, facilitating more **efficient topological routing**
 - **multi-layer management** for handling configuration, performance and security.
- **Demonstrate** the applicability and benefits of this approach and its built-in functions in **three use-cases**
 - **Datacenter, Distributed cloud, Carrier network**

PRISTINE – (SW) Outcomes

- **PRISTINE will be:**
 - **Developing a Software Development Kit (SDK):**
 - From the IRATI prototype sources-base
 - Enable to customize the behaviour of the DIFs
 - Allow to plug-in/out policies dynamically
 - Policies in the transport layer
 - Policies in the management layers
 - **Developing innovative policies**
 - New policies, defined in the project will be: developed, tested and integrated
 - **Developing the first DIF Management System (DMS)**
 - Manages the DIFs in multiple Processing Systems (same administrative domain)
 - **IPC Process and IPC Manager daemons enhancements**
 - Management Agent
 - Coordinates the loading/installation/removal/unloading of policies in the Processing System
 - Reliability aspects
 - Measurements
 - Performances
 - ...

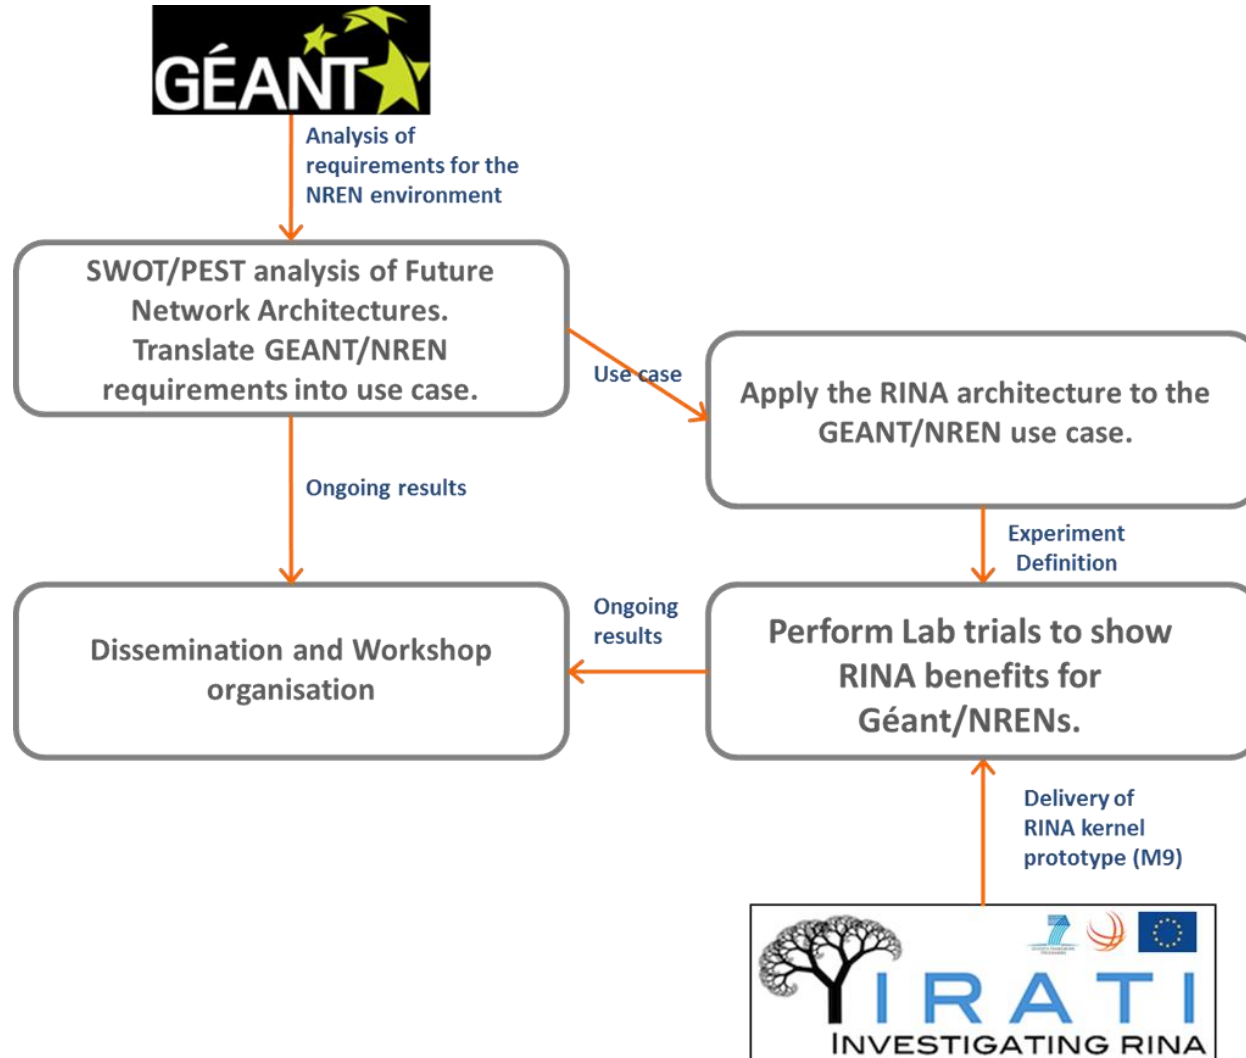
GEANT3+ IRINA – OVERVIEW

- Investigating **RINA** as the next generation GEANT and **NREN** network **architecture** (IRINA)
- GEANT3+ project
 - Starts **Oct 2013**, ends **March 2015** (18 months)
- 4 Partners:
 - [**Research**] iMinds VZW (Belgium)
 - [**Research**] Fundació Privada i2CAT (Spain)
 - [**Research**] Waterford Institute of Technology – Telecommunications Software & Systems Group (Ireland)
 - [**SME**] Nextworks s.r.l. (Italy)

IRINA - Objectives

- Proposes to study RINA as the foundation of the next generation NREN and GEANT network architectures.
- Targets the following goals:
 - **Make a comparative study of RINA vs. the current networking SoTA** and the most relevant clean-slate architectures under research.
 - **Perform a use-case study** of how RINA could be better used in the NREN scenario
 - considering different deployment options, and illustrating the benefits that RINA can bring in terms of multi-homing, mobility, quality of service, programmability, virtualization and network management.
 - **Showcase a lab trial of the use-case study**
 - Utilizing a customized version of the FP7 IRATI stack, and the experimental facilities contributed by the project partners.
 - **Involve the NREN and GEANT community** in the different steps of the project, in order to discuss the project approach, the findings and to get valuable feedback.
 - The organization of a network architectures workshop in cooperation with GN3+ JRA1 will be a key instrument to achieve this objective.

IRINA – Overview/Objectives



STANDARDISATION

- **RINA addresses concerns identified by FN:**

“Even though the current Internet is such an essential infrastructure, we see that there are many concerns about the following technical aspects of the current Internet, including IP based networks: **scalability**, ubiquity, **security**, **robustness**, **mobility**, **heterogeneity**, **Quality of Service (QoS)**, **re-configurability**, context-awareness, **manageability**, economics, etc.”

- **RINA IPC Specification Reference Model**
 - Basic concept of distributed systems
 - Distributed applications
 - **Distributed InterProcess Communication (IPC API)**
- **DAF Base Specifications**
 - **Common Application Establishment Phase**
 - **Common Distributed Application Protocol**
 - **IPC Resource Manager Specification**
 - **DIF Allocator Specification**

- **DIF Base Specifications**

- Data Transfer Service Definition
- Specification template for a Generic DIF Delimiting module
- Error and Flow Control Protocol Specification (DTP + DTCP)
- Relaying and Multiplexing Task Specification
- Specification Template for a DIF SDU Protection Module
- Specification Template for a Generic DIF SDU Protection Module
- Basic Enrollment Specification
- Flow Allocator Specification

- **Policy Specifications**

- CRC 16 SDU protection module
- DIF HDLC like SDU protection
- DIF TCP UDP like SDU protection module
- Retransmission timer expiry policy for a TCP like DIF
- Round Trip Time (RTT) estimator policy for a TCP like DIF
- Delimiting module for operating over the public Internet
- Delimiting module for demo DIF

- **Pro-forma for policy specifications**

- **No need to standardise everything**

- **Shim-eth-vlan**
- **Shim-tcp-udp**
- **Shim-hv**

Upcoming workshops

- Globecom Workshop “Alternatives to TCP/IP”
 - 8-12 December, Austin TX US
- RINA workshop
 - 28-29 January 2015, Ghent Belgium
- TERENA TNC 2015
 - 15-18 June, Porto, Portugal
- Summer school ~ August 2015